

# Fliprobo Assignment 2 Machine learning (assignment 5)

Q1. R-squared or Residual Sum of Squares (RSS) which one of these two is a better measure of goodness of fit model in regression and why?

Answer : Comparison and Suitability

1. Scale and Interpretability:

RSS: The value of RSS is dependent on the scale of the data. It is not standardized, which makes it difficult to compare across different datasets or models.

R-squared: It is a standardized measure (ranging from 0 to 1), making it easier to interpret and compare across different datasets and models.

2. Explained Variance:

RSS: While it provides a measure of how well the model fits the data, it doesn't provide information on the proportion of the total variance explained by the model.

R-squared: Directly indicates the proportion of variance explained, providing a clear understanding of the model's explanatory power.

3. Context of Use:

RSS: Useful when comparing models with the same dataset to understand which model has smaller residuals.

R-squared: More commonly used in reporting the goodness of fit because it provides a clear and intuitive measure of how well the independent variables explain the variance in the dependent variable.

R-squared is generally considered a better measure of the goodness of fit for regression models compared to RSS because:

It is standardized, making it easier to interpret and compare.

It directly indicates the proportion of the variance explained by the model, providing a clear measure of the model's explanatory power.

It is widely used and recognized, making it a more accepted metric in reporting and interpreting regression results.

Therefore, while RSS is valuable in specific contexts, R-squared is typically the preferred measure for assessing the goodness of fit in regression analysis.

Q2. What are TSS (Total Sum of Squares), ESS (Explained Sum of Squares) and RSS (Residual Sum of Squares) in regression. Also mention the equation relating these three metrics with each other.

Answer: TSS measures the total variance in the dependent variable  $y$ . It quantifies the total variability in the response variable. TSS represents the total variability in the data.

ESS measures the amount of variance explained by the regression model. It quantifies how much of the total variability in  $y$  is explained by the fitted regression model. ESS represents the variability explained by the regression model.

RSS measures the amount of variance that is not explained by the regression model. It quantifies the variability in  $y$  that remains after fitting the model. RSS represents the variability that the model fails to explain.

Relationship :  $TSS = ESS + RSS$

Q3. What is the need of regularization in machine learning?

Answer : Regularization is a crucial concept in machine learning, especially in the context of regression and classification models. The primary need for regularization arises from the challenges of overfitting and underfitting, which impact the model's ability to generalize well to unseen data.

1. Prevent Overfitting:

Overfitting occurs when a model learns not only the underlying pattern in the training data but also the noise and outliers. This leads to a model that performs well on training data but poorly on new, unseen data.

Regularization techniques add a penalty to the loss function for large coefficients, effectively discouraging the model from becoming too complex and sensitive to the noise in the training data.

2. Improve Generalization:

Generalization refers to the model's ability to perform well on new, unseen data. Regularization helps improve generalization by simplifying the model, ensuring it captures the underlying trend rather than the specific details of the training data.

3. Reduce Variance:

High-variance models are sensitive to the fluctuations in the training data, leading to different models if trained on different subsets of the data. Regularization reduces the variance of the model by making it less sensitive to small changes in the training data.

Q4. What is Gini-impurity index?

Answer : Gini impurity index is a metric used to evaluate the quality of a split in decision trees, particularly in classification problems. It measures the degree of impurity or disorder in a dataset, with the goal of creating the most homogeneous nodes possible in the decision tree. The Gini impurity ranges between 0 and 0.5 for binary classification. A Gini impurity of 0 indicates a pure node, where all the items belong to a single class. The maximum impurity, 0.5, occurs when the items are evenly split between two classes. Lower Gini impurity values indicate purer nodes, which means a more homogeneous set of items.

The Gini impurity index is a vital concept in decision tree algorithms, helping to create splits that result in the most homogeneous child nodes. By minimizing the Gini impurity at each split, decision trees aim to improve their predictive accuracy and interpretability.

Q5. Are unregularized decision-trees prone to overfitting? If yes, why?

Answer : Yes, unregularized decision trees are prone to overfitting. This tendency arises due to several inherent characteristics of decision trees:

Complexity and Flexibility:

Decision trees are highly flexible and can create very complex models that perfectly fit the training data, including noise and outliers.

This flexibility allows the tree to split the data into increasingly smaller subsets until each leaf node contains only one instance or instances with the same label, leading to a model that captures the idiosyncrasies of the training data rather than the underlying patterns.

#### Deep Trees:

Without any constraints, decision trees can grow very deep, resulting in a large number of nodes and branches.

Deep trees with many splits tend to have very small subsets of data in their leaf nodes, which can cause the model to be overly sensitive to minor variations in the training data.

#### High Variance:

Unregularized decision trees often exhibit high variance, meaning they perform well on the training data but poorly on new, unseen data (i.e., they do not generalize well).

This high variance is a direct consequence of the model's complexity and its ability to fit the training data too closely.

#### Q6. What is an ensemble technique in machine learning?

Answer : Ensemble techniques in machine learning involve combining multiple models to improve overall performance, reduce variance, and enhance robustness. These techniques leverage the strengths of various individual models to produce more accurate and reliable predictions than any single model could achieve alone. The basic idea is that by aggregating the predictions of multiple models, the ensemble can correct for the weaknesses and errors of the individual models.

#### Types of ensemble techniques

1. Bootstrap aggregating (bagging)
2. Boosting
3. Voting
4. Stacking (Stacked generalization)

Q7. What is the difference between Bagging and Boosting techniques?

Answer : Both bagging and boosting are powerful ensemble techniques that improve the performance of machine learning models, but they do so in different ways. Bagging focuses on reducing variance by training multiple models in parallel on different subsets of the data, while boosting focuses on reducing bias and variance by training models sequentially, each new model correcting the errors of the previous ones. The choice between bagging and boosting depends on the specific problem and the nature of the data.

Q8. What is out-of-bag error in random forests?

Answer :

The out-of-bag error in random forests provides a reliable and efficient way to estimate the model's prediction error without needing a separate validation set. It leverages the inherent bootstrap sampling of random forests, making full use of the dataset and providing an unbiased error estimate that helps in assessing the model's generalization performance..Out-of-bag (OOB) error is a method used to estimate the prediction error of random forests without the need for a separate validation set. It leverages the bootstrap sampling technique inherent in the construction of random forests.

Q9. What is K-fold cross-validation?

Answer : K-fold cross-validation is a robust and widely used technique for evaluating the performance of a machine learning model. It helps in assessing how the model will generalize to an independent dataset, which is crucial for preventing overfitting and ensuring that the model performs well on unseen data.

Definition and Process

K-fold cross-validation involves splitting the dataset into K equally-sized (or nearly equally-sized) subsets, known as "folds." The process is as follows:

Divide the Data: Split the entire dataset into K folds. Each fold contains approximately the same number of instances, ensuring that the distribution of the target variable is consistent across folds.

Train and Validate: For each fold:

Use  $K-1$  folds for training the model.

Use the remaining 1 fold for validating the model.

This process is repeated K times, each time using a different fold as the validation set and the remaining K-1 folds as the training set.

Compute Performance Metrics: After the model is trained and evaluated on each of the K validation sets, aggregate the performance metrics (e.g., accuracy, precision, recall, F1 score) across all folds to obtain an overall assessment of the model's performance.

Q10. What is hyper parameter tuning in machine learning and why it is done?

Answer : Hyperparameter tuning is the process of optimizing the hyperparameters of a machine learning model to improve its performance on a given task. Hyperparameters are settings or configurations that are external to the model and whose values cannot be estimated from the training data. They need to be set prior to training the model and can significantly influence the model's performance.

Hyperparameter tuning is a critical step in the machine learning workflow that can significantly impact model performance. By carefully selecting and optimizing hyperparameters, practitioners can enhance model accuracy, generalization, and efficiency, leading to better predictive performance on new data.

Hyperparameter tuning is crucial because the choice of hyperparameters can have a significant impact on the performance of a machine learning model. Properly tuned hyperparameters can:

Improve Model Performance: Optimizing hyperparameters can lead to higher accuracy, precision, recall, F1 score, or other relevant metrics.

Enhance Generalization: Well-tuned hyperparameters help the model generalize better to new, unseen data, reducing overfitting.

**Optimize Training Time:** Efficient hyperparameters can lead to faster convergence and reduce the computational cost of training the model.

**Balance Bias and Variance:** Proper tuning can help in finding the right trade-off between bias and variance, which is essential for good model performance.

Q11. What issues can occur if we have a large learning rate in Gradient Descent?

Answer : A large learning rate in gradient descent can lead to several issues that negatively impact the training process and the performance of a machine learning model. Here are the primary problems associated with having a large learning rate:

### Issues with Large Learning Rate

#### 1. Divergence:

Instead of converging to the minimum of the loss function, the model's parameters may start oscillating or even diverging, moving away from the minimum.

#### 2. Oscillation:

The model parameters may jump back and forth around the minimum, failing to settle down.

#### 3. Instability:

The training process becomes unstable, with the loss function fluctuating wildly.

#### 4. Poor Convergence:

The model may take longer to converge or may not converge at all.

#### 5. Suboptimal Solutions:

The model may converge to a suboptimal solution or local minimum that is not the best possible solution.

Q12. Can we use Logistic Regression for classification of Non-Linear Data? If not, why?

Answer : Logistic Regression is primarily a linear classifier, meaning it works well for datasets where the classes can be separated by a linear decision boundary. When dealing with non-linear data, where the relationship between features and the target class is not linearly separable, Logistic Regression might not perform well.

Why Logistic Regression Struggles with Non-Linear Data

1. Linear Decision Boundary:

Nature: Logistic Regression models a linear decision boundary between classes. This means it tries to find a straight line (or hyperplane in higher dimensions) that separates the classes.

Limitation: If the true boundary between classes is non-linear, Logistic Regression will fail to capture this complexity, leading to poor classification performance.

2. Feature Space:

Linear Separation: Logistic Regression assumes that classes are linearly separable in the feature space. For non-linearly separable data, this assumption does not hold.

Q13. Differentiate between Adaboost and Gradient Boosting.

Answer :

Adaboost focuses on improving the performance of weak learners by adjusting the weights of incorrectly classified instances. It iteratively trains weak learners, giving more attention to the instances that previous learners misclassified. This process helps to create a strong classifier from a series of weak classifiers.



### *Algorithm*

Initial Weights: Initially, all training instances are assigned equal weights.

Iteration: In each iteration, a weak learner is trained on the weighted training data. The error rate of the learner is then calculated. Weights of misclassified instances are increased, making them more influential in the next iteration.

Final Model: The final prediction is a weighted majority vote (for classification) or a weighted sum (for regression) of all the weak learners' predictions.

Adaboost focuses on instances that were misclassified by previous weak learners. By increasing the weights of these difficult instances, subsequent learners are more likely to correct these errors, thereby improving the overall model performance.

The weight of each training instance is updated based on its classification correctness. Misclassified instances receive higher weights, while correctly classified instances receive lower weights. This dynamic adjustment helps the model to focus on harder-to-classify instances.

The final model combines the weak learners using a weighted majority vote (in classification) or a weighted sum (in regression). Each learner's contribution is proportional to its accuracy, giving more weight to more accurate learners.

VS

Gradient Boosting builds an additive model in a forward stage-wise manner by optimizing a loss function through gradient descent. Each new model is trained to correct the residual errors of the combined ensemble of all previous models.

### *Algorithm*

Initial Model: Start with an initial prediction, such as the mean of the target values in regression.

Iteration: In each iteration, compute the residuals (errors) of the current model. Train a new weak learner to predict these residuals and add this learner to the ensemble. The contribution of each learner is scaled by a learning rate.

Final Model: The final prediction is the sum of the initial model and all subsequent weak learners, each weighted by the learning rate.

Each new weak learner is trained to predict the residuals (errors) of the ensemble of previous learners. This approach effectively performs gradient descent on the loss function, iteratively reducing the overall error.

Instead of adjusting instance weights, Gradient Boosting updates the predictions by adding new models that correct the residuals of the previous models. This additive correction helps in progressively improving the model.

Models are combined by summing their contributions, often scaled by a learning rate. This creates an additive model where each new learner incrementally improves the overall prediction accuracy.

## Key Differences

### *Weight Update*

Adaboost adjusts the weights of the training instances, giving more weight to misclassified instances to focus subsequent learners on these harder cases. Gradient Boosting, on the other hand, adjusts the model predictions by adding new models that target the residual errors of the previous models.

### *Focus*

Adaboost increases focus on misclassified instances by changing their weights, making future learners pay more attention to these instances. Gradient Boosting focuses on minimizing the residual errors from previous learners, thereby iteratively improving the model's predictions.

### *Loss Function*

Adaboost implicitly minimizes the exponential loss function by adjusting instance weights based on classification accuracy. Gradient Boosting explicitly minimizes a specified loss function (e.g., mean squared error for regression) through gradient descent.

### *Combination of Models*

In Adaboost, models are combined using a weighted majority vote or sum, with each model's weight based on its accuracy. In Gradient Boosting, models are combined in an

additive manner, where each new model's contribution is scaled by a learning rate to incrementally improve the prediction.

### *Iteration Dependency*

In Adaboost, the weights of instances depend on their classification correctness in the previous iteration. In Gradient Boosting, the residuals depend on the errors of the combined model from all previous iterations, with each new model correcting these residuals.

### *Robustness to Noise*

Adaboost can be less robust to noisy data because it increases the weights of hard-to-classify instances, which may include noise. Gradient Boosting is generally more robust to overfitting, especially with regularization techniques, but can still overfit without proper tuning.

### *Learning Rate*

Adaboost typically does not use a learning rate, whereas Gradient Boosting uses a learning rate to control the contribution of each model, which helps in preventing overfitting and improving generalization

Q14. What is bias-variance trade off in machine learning?

Answer : The bias-variance trade-off is a fundamental concept in machine learning that describes the balance between two sources of error that affect the performance of predictive models: bias and variance. Understanding and managing this trade-off is crucial for building models that generalize well to new, unseen data.

### **Bias**

Bias refers to the error introduced by approximating a real-world problem, which may be complex, by a simplified model. High bias means the model makes strong assumptions about the data, leading to systematic errors in predictions.

High Bias: Occurs when the model is too simple to capture the underlying patterns in the data, resulting in underfitting. Examples include linear models applied to non-linear data.

Impact: High bias models tend to have high training error and high validation error, as they cannot capture the complexities of the data.

## Variance

Variance refers to the error introduced by the model's sensitivity to small fluctuations in the training data. High variance means the model is overly complex and fits the noise in the training data rather than the actual pattern.

High Variance: Occurs when the model is too complex, leading to overfitting.

Examples include deep decision trees or models with too many parameters relative to the number of observations.

Impact: High variance models tend to have low training error but high validation error, as they perform well on training data but poorly on new, unseen data.

## Trade-Off

The bias-variance trade-off is the balance between these two sources of error:

Low Bias and High Variance: The model overfits the training data, capturing noise and leading to poor generalization to new data.

High Bias and Low Variance: The model underfits the training data, failing to capture important patterns and leading to poor performance on both training and validation data.

Optimal Trade-Off: The goal is to find a balance where both bias and variance are minimized, resulting in a model that generalizes well to new data.

## Visualization

High Bias (Underfitting): The model is too simple, resulting in a high error on both training and test data.

High Variance (Overfitting): The model is too complex, resulting in low training error but high test error.

Optimal Model: Achieves a balance where the training error and test error are both low.

Q15. Give short description each of Linear, RBF, Polynomial kernels used in SVM.

Answer : Linear Kernel

The linear kernel is the simplest kernel function used in SVM.

It computes the dot product of two feature vectors and adds a constant, essentially representing a linear relationship between the features.

RBF (Radial Basis Function) Kernel

The RBF kernel, also known as the Gaussian kernel, is a popular choice for non-linear data.

It measures the similarity between two data points based on their distance, using a Gaussian function.

Polynomial Kernel

The polynomial kernel represents the similarity of vectors in a feature space over polynomials of the original variables, allowing for non-linear relationships.